

Cross-lens 2×2 — reliance × training-data overlap
3,580 spans

context-driven (CTI ≥ 0.30)	949 (26.5%)	2,329 (65.1%)
mixed	78 (2.2%)	180 (5.0%)
parametric (CTI ≤ 0.05)	5 (0.1%)	39 (1.1%)
	pretraining match	no match

dashed cell: parametric and untraced-in-pretraining (“no visible source”)